

# R Basics

Tomasz Mądry

2025-07-09

# Overview

- ▶ Important R libraries
  - ▶ dplyr
  - ▶ ggplot2 + scales
  - ▶ stringr
- ▶ Exercises

# dplyr - Introduction

## Overview

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

```
mutate() # adds new variables that are functions of existing variables
select() # picks variables based on their names.
filter() # picks cases based on their values.
summarise() # reduces multiple values down to a single summary.
arrange() # changes the ordering of the rows.
```

These all combine naturally with `group_by()` which allows you to perform any operation “by group”.

## More info

<https://dplyr.tidyverse.org/>

## dplyr - Data set preview

Load exemplary data set:

```
data(mtcars)
```

Preview the data:

```
head(mtcars, 6) # display first 6 rows
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

## dplyr - Data set preview

```
tail(mtcars, 2) # display last 2 rows
```

```
##               mpg cyl  disp  hp  drat   wt  qsec vs  am gear carb
## Maserati Bora 15.0   8   301 335  3.54  3.57 14.6  0   1    5    8
## Volvo 142E    21.4   4   121 109  4.11  2.78 18.6  1   1    4    2
```

```
nrow(mtcars) # display number of rows
```

```
## [1] 32
```

```
ncol(mtcars) # display number of cols
```

```
## [1] 11
```

## dplyr - Pipe %>%

Pipe (ctrl + shift + M):

```
%>% # magrittr pipe
```

Output from the left function is passed as input to right function:

```
mtcars %>% head()
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

## dplyr - Subset rows: filter()

Select rows meeting certain conditions.

Example: Cars with 6 cylinders

```
mtcars %>% filter(cyl == 6) %>% head(6)
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4

## dplyr - Choose columns: select()

Pick specific columns from the dataset.

Example: Select mpg, cyl, and hp columns

```
mtcars %>% select(mpg, cyl, hp) %>% head(6)
```

##	mpg	cyl	hp
## Mazda RX4	21.0	6	110
## Mazda RX4 Wag	21.0	6	110
## Datsun 710	22.8	4	93
## Hornet 4 Drive	21.4	6	110
## Hornet Sportabout	18.7	8	175
## Valiant	18.1	6	105



## dplyr - Add/modify columns: mutate()

Create new columns or modify existing ones.

Example: Calculate fuel consumption in liters per 100 km

```
mtcars %>%  
  select(mpg, cyl) %>%  
  mutate("l per 100km" = round((235.215 / mpg), 0)) %>%  
  head(6)
```

##	mpg	cyl	l per 100km
## Mazda RX4	21.0	6	11
## Mazda RX4 Wag	21.0	6	11
## Datsun 710	22.8	4	10
## Hornet 4 Drive	21.4	6	11
## Hornet Sportabout	18.7	8	13
## Valiant	18.1	6	13

## dplyr - Sort rows: arrange()

Order rows by column values ascending or descending.

Example: Sort cars by mpg descending

```
mtcars %>%  
  arrange(desc(mpg)) %>%  
  head(6)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2

## dplyr - Aggregate data by groups: group\_by() + summarize()

Group data and compute summary statistics per group.

Example: Average mpg per cylinder count

```
mtcars %>%  
  group_by(cyl) %>%  
  summarize(avg_mpg = mean(mpg),  
            count = n())
```

```
## # A tibble: 3 x 3  
##   cyl avg_mpg count  
##   <dbl> <dbl> <int>  
## 1     4   26.7    11  
## 2     6   19.7     7  
## 3     8   15.1    14
```

## dplyr - Alternative for group\_by() + summarize()

Grouping can be done within summarize() function.

Example: Average mpg per cylinder count

```
mtcars %>%  
  summarize(.by = cyl,  
            avg_mpg = mean(mpg),  
            count = n())
```

```
##   cyl  avg_mpg count  
## 1    6 19.74286     7  
## 2    4 26.66364    11  
## 3    8 15.10000    14
```

## dplyr - Remove duplicate rows `distinct()`

Returns unique rows from a data frame.

Example: Get unique combinations of cylinders and gears:

```
mtcars %>% select(cyl, gear) %>% distinct()
```

##	cyl	gear
## Mazda RX4	6	4
## Datsun 710	4	4
## Hornet 4 Drive	6	3
## Hornet Sportabout	8	3
## Toyota Corona	4	3
## Porsche 914-2	4	5
## Ford Pantera L	8	5
## Ferrari Dino	6	5

## dplyr - Rename columns rename()

Changes column names in a data frame.

Example: Rename mpg to miles\_per\_gallon

```
mtcars %>%  
  select(mpg, cyl, disp, hp) %>%  
  rename(miles_per_gallon = mpg) %>%  
  head(6)
```

	miles_per_gallon	cyl	disp	hp
## Mazda RX4	21.0	6	160	110
## Mazda RX4 Wag	21.0	6	160	110
## Datsun 710	22.8	4	108	93
## Hornet 4 Drive	21.4	6	258	110
## Hornet Sportabout	18.7	8	360	175
## Valiant	18.1	6	225	105

## dplyr - Rename columns pull()

Useful for quickly extracting a single column.

Example: Get the mpg column as a vector

```
mtcars %>%  
  head(6) %>%  
  pull(mpg)
```

```
## [1] 21.0 21.0 22.8 21.4 18.7 18.1
```

## dplyr - Perform operations row-by-row rowwise()

Applies functions across columns for each row.

Example: Calculate row-wise average of mpg, hp, and wt

```
mtcars %>%  
  head(2) %>%  
  select(mpg, hp, wt) %>%  
  rowwise() %>%  
  mutate(avg_value = mean(c(mpg, hp, wt)))
```

```
## # A tibble: 2 x 4  
## # Rowwise:  
##   mpg    hp    wt avg_value  
##   <dbl> <dbl> <dbl>   <dbl>  
## 1    21   110  2.62    44.5  
## 2    21   110  2.88    44.6
```



## dplyr - Conditional column creation if\_else()

Vectorized conditional logic (used inside mutate()).

Example: Label cars as “high” or “low” mpg

```
mtcars %>%  
  tail(4) %>%  
  select(mpg, hp) %>%  
  mutate(mpg_type = if_else(mpg > 20, "high", "low"))
```

		mpg	hp	mpg_type
##	Ford Pantera L	15.8	264	low
##	Ferrari Dino	19.7	175	low
##	Maserati Bora	15.0	335	low
##	Volvo 142E	21.4	109	high

# ggplot2 - Introduction

ggplot2 is a powerful and flexible R package for creating declarative, layered graphics. It is based on the Grammar of Graphics, which allows users to build plots step-by-step by adding components.

Key Features:

- ▶ Intuitive syntax using + to add layers
- ▶ Supports a wide range of plots: scatter, bar, line, boxplot, histogram, etc.
- ▶ Enables layered plotting (data + geoms + stats + scales + themes)
- ▶ Powerful faceting system for multi-panel plots
- ▶ Highly customizable aesthetics (color, size, shape, labels, etc.)

## ggplot2 - Core functions

- ▶ `ggplot()` – creates a plotting object
- ▶ `geom_*()` – add geometric objects (points, bars, lines)
- ▶ `aes()` – define aesthetic mappings (e.g., x, y, color)
- ▶ `facet_wrap()` / `facet_grid()` – split plots by variables
- ▶ `scale_*()` – control axes, colors, and legends
- ▶ `theme_*()` – customize non-data elements (background, font)

### More info

<https://ggplot2.tidyverse.org/>

## scales - Introduction

scales is an R package that provides tools to customize how data values are mapped to visual properties (like axes, colors, and labels) in plots—especially those created with ggplot2.

Key Features:

- ▶ Format numbers and dates for axis labels (e.g., currency, percentages, commas)
- ▶ Create and customize color palettes and gradients
- ▶ Apply transformations (log, reverse, square root) to scales
- ▶ Generate breaks and labels automatically or manually for axes
- ▶ Support for continuous, discrete, and date/time scales

## scales - Core functions

- ▶ Label functions: `label_dollar()`, `label_percent()`, `label_comma()`, etc. — format axis and legend labels
- ▶ Color scales: `scale_color_gradient()`, `scale_fill_brewer()`, etc. — define color mappings
- ▶ Transformations: `trans_new()`, built-in like `log_trans()`, `reverse_trans()` — modify scale behavior
- ▶ Breaks and limits: control where ticks and labels appear on axes

### More info

<https://scales.r-lib.org/>

## ggplot2 + scales: diamonds data set

A data set containing the prices and other attributes of almost 54,000 diamonds. Part of the ggplot2 package.

Key columns:


- ▶ price: price in US dollars
- ▶ carat: weight of the diamond
- ▶ cut, color, clarity: quality metrics
- ▶ x, y, z: dimensions

```
## # A tibble: 3 x 10
```

```
##   carat cut      color clarity depth table price      x      y      z
##   <dbl> <ord>   <ord> <ord>   <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1  0.23 Ideal    E     SI2     61.5    55   326  3.95  3.98  2.43
## 2  0.21 Premium E     SI1     59.8    61   326  3.89  3.84  2.31
## 3  0.23 Good    E     VS1     56.9    65   327  4.05  4.07  2.31
```

## ggplot2 + scales: plot base

```
ggplot(data = diamonds, aes(x = carat, y = price))
```



## ggplot2 + scales: geoms


Add layers to represent data.

- ▶ `geom_point()` – scatterplot
- ▶ `geom_bar()`, `geom_col()` – bar charts
- ▶ `geom_histogram()` – histogram
- ▶ `geom_boxplot()` – boxplot



ggplot2 + scales: geom\_point()

```
ggplot(data = diamonds, aes(x = carat, y = price)) + geom_point()
```




## ggplot2 + scales: Stats (stat\_\*)

Statistical transformations.

- ▶ `stat_smooth(method = "lm")` – regression line


```
ggplot(diamonds, aes(carat, price)) +  
  geom_point(alpha = 0.1) +  
  stat_smooth(formula = y ~ x, method = "lm", col = "red") +  
  theme_minimal() # select nice theme here
```



## ggplot2 + scales: Themes

Control non-data visuals (font, background, gridlines).


```
ggplot(diamonds, aes(carat, price)) +  
  geom_point(alpha = 0.1) +  
  stat_smooth(formula = y ~ x, method = "lm", col = "red")
```



## ggplot2 + scales: Themes

Control non-data visuals (font, background, gridlines).


```
ggplot(diamonds, aes(carat, price)) +  
  geom_point(alpha = 0.1) +  
  stat_smooth(formula = y ~ x, method = "lm", col = "red") +  
  theme_bw()
```



## ggplot2 + scales: Themes

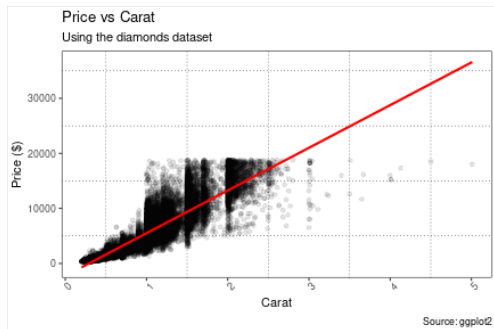
Control non-data visuals (font, background, gridlines).

```
ggplot(diamonds, aes(carat, price)) +  
  geom_point(alpha = 0.1) +  
  stat_smooth(formula = y ~ x, method = "lm", col = "red") +  
  theme_bw() +  
  theme(plot.background = element_rect(fill = "transparent"),  
        axis.text.x = element_text(angle = 45),  
        panel.grid.major = element_blank(),  
        panel.grid.minor = element_line(linetype = 3, color = "black"))
```



## ggplot2 + scales: Labels and Titles

```
labs(title = "Price vs Carat",  
      subtitle = "Using the diamonds dataset",  
      x = "Carat", y = "Price ($)",  
      caption = "Source: ggplot2")
```




## ggplot2 + scales: Scales

Control axes and legends.

```
xlab("Weight (carats)") +  
ylab("Price") +  
scale_x_continuous(breaks = seq(0, 3, 0.5))  
scale_y_continuous(labels = scales::label_dollar()) # or simply: scales::dollar
```

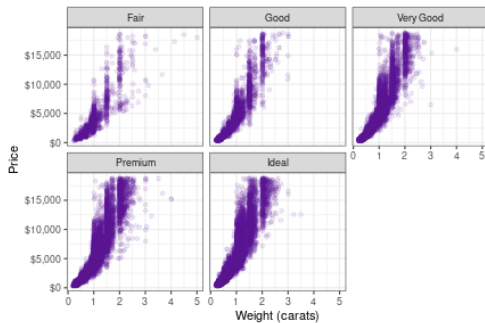
```
## Warning: Removed 32 rows containing missing values or values outside the scale range for 'Weight (carats)'.  
## ('geom_point()').
```



## ggplot2 + scales: Facets

Split data into subplots.

► `facet_wrap(~cut)`

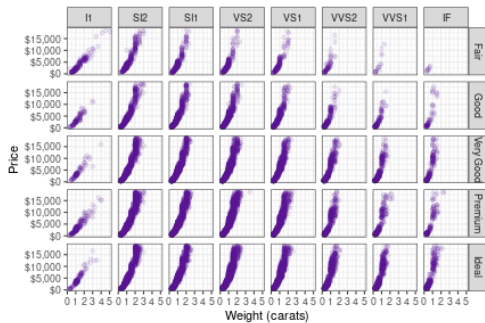




## ggplot2 + scales: Facets

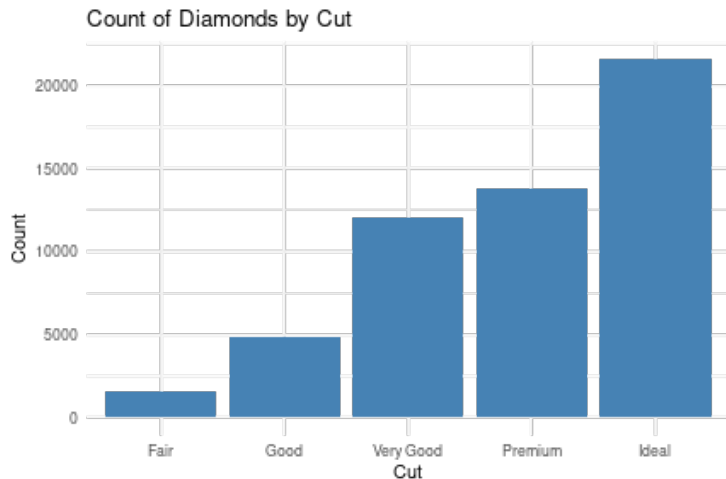
Split data into subplots.

► `facet_grid(cut ~ clarity)`




## Additional geometrics: geom\_bar

Count of diamonds by cut.



## Additional geometrics: `geom_line`

Average price by carat.



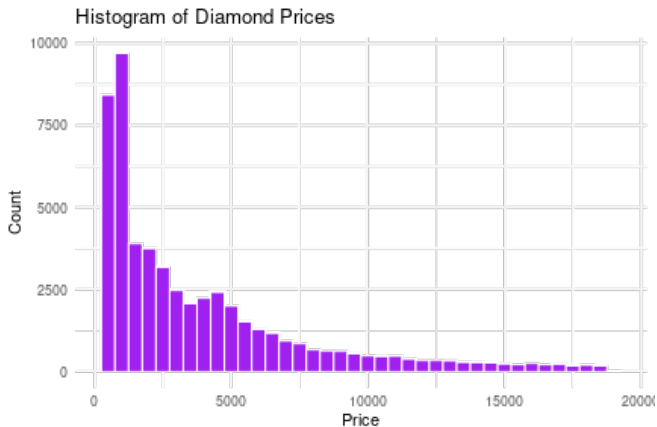
## Additional geometrics: `geom_boxplot`

Price distribution by cut.



## Additional geometrics: `geom_histogram`

Distribution of diamond prices.



## stringr - Introduction

stringr is an R package designed to make working with strings (character data) easier, consistent, and more user-friendly. It provides a cohesive set of functions that simplify common string manipulation tasks such as detecting, extracting, replacing, and splitting strings.

Function	Purpose
<code>str_length()</code>	Returns the number of characters in a string
<code>str_sub()</code>	Extracts substrings based on position
<code>str_detect()</code>	Checks if a pattern exists in a string
<code>str_subset()</code>	Extracts strings that match a pattern
<code>str_replace()</code>	Replaces the first match of a pattern
<code>str_replace_all()</code>	Replaces all matches of a pattern
<code>str_split()</code>	Splits strings into pieces based on a pattern
<code>str_trim()</code>	Removes leading and trailing whitespace

## stringr - Examples

```
library(stringr)
```

```
text <- "The quick brown fox jumps over the lazy dog."
```

```
# Detect if 'fox' is in the string  
str_detect(text, "fox")
```

```
## [1] TRUE
```

```
# Extract the first 9 characters  
str_sub(text, 1, 9)
```

```
## [1] "The quick"
```

```
# Replace 'fox' with 'cat'  
str_replace(text, "fox", "cat")
```

```
## [1] "The quick brown cat jumps over the lazy dog."
```

## stringr in real-life: modify gene\_ids

```
df <- tibble(gene_id = c("ENSG000000000457.14", "ENSG000000000005.6", "ENSG00000310545.1"),  
             biotype = c("protein_coding", "protein_coding", "lncRNA"))  
df
```

```
## # A tibble: 3 x 2  
##   gene_id          biotype  
##   <chr>           <chr>  
## 1 ENSG000000000457.14 protein_coding  
## 2 ENSG000000000005.6 protein_coding  
## 3 ENSG00000310545.1  lncRNA
```

```
df %>% mutate(gene_id = str_remove(gene_id, "\\\\.\\d+$"))
```

```
## # A tibble: 3 x 2  
##   gene_id          biotype  
##   <chr>           <chr>  
## 1 ENSG000000000457 protein_coding  
## 2 ENSG000000000005 protein_coding  
## 3 ENSG00000310545  lncRNA
```



## stringr in real-life: modify biotypes

```
df
```

```
## # A tibble: 3 x 2
##   gene_id          biotype
##   <chr>          <chr>
## 1 ENSG00000000457.14 protein_coding
## 2 ENSG00000000005.6  protein_coding
## 3 ENSG00000310545.1  lncRNA
```

```
df %>% mutate(biotype = str_replace(biotype, "protein_coding", "pc"))
```

```
## # A tibble: 3 x 2
##   gene_id          biotype
##   <chr>          <chr>
## 1 ENSG00000000457.14 pc
## 2 ENSG00000000005.6  pc
## 3 ENSG00000310545.1  lncRNA
```

## Excercises

## Exercise 1:

### iris Dataset: Average Sepal Length by Species

**Dataset:** iris

**Task:**

- ▶ Use `dplyr` to compute the average sepal length per species.
- ▶ Use `ggplot2` to create a bar plot of the results.

**Hint:** Use `geom_bar(stat = "identity")` or `geom_col()`!

## Excercise 2:

### `mtcars`: Horsepower vs MPG by Cylinders

**Dataset:** `mtcars`

**Task:**

- ▶ Use `dplyr` to group by the number of cylinders (`cyl`) and compute the average `mpg` and `hp`.
- ▶ Create a scatter plot of avg horsepower vs avg mpg, colored by number of cylinders.

**Hint:** Use `factor(cyl)` for color aesthetics.

## Excercise 3:

### ToothGrowth: Effect of Dose on Tooth Length

**Dataset:** ToothGrowth

**Task:**

- ▶ Use `dplyr` to calculate mean tooth length for each dose.
- ▶ Create a line plot showing `tooth length` vs `dose`, colored by supplement type (`supp`).

**Hint:** Don't forget to `group_by` both: `dose` and `supp`!

## Excercise 4:

### C02: Uptake Rate by Treatment

**Dataset:** C02

**Task:**

- ▶ Filter for Type == "Quebec".
- ▶ Use dplyr to get average uptake per treatment and concentration.
- ▶ Use ggplot2 to plot a line graph of uptake vs concentration for each treatment.

## Excercise 5:

### ChickWeight: Growth Over Time

**Dataset:** ChickWeight

**Task:**

- ▶ Use `dplyr` to filter for a few selected chicks (e.g., `Chick %in% c(1, 5, 10)`).
- ▶ Use `ggplot2` to create a line plot of weight over time per chick.

## Excercise 6:

### mtcars: Boxplot of MPG by Transmission

**Dataset:** mtcars

**Task:**

- ▶ Use `dplyr` to mutate `am` to a factor labeled "Automatic" and "Manual".
- ▶ Create a boxplot of `mpg` by transmission type (`am`).



## Excercise 7:

### iris: Sepal Width Distribution

**Dataset:** iris

**Task:**

- ▶ Use `ggplot2` to make a histogram of `sepal width`, filled by `species`.
- ▶ Use `facet_wrap()` to show separate plots per `species`.

## Excercise 8:

### USArrests: Violent Crime by Region

**Dataset:** USArrests

**Task:**

- ▶ Add a column for region (use `state.region` from base R, it matches row names of `USArrests`).
- ▶ Use `dplyr` to get average murder and assault rate per region.
- ▶ Create a bar plot of average assault rate per region.

**Hint:** To match regions use: `state.region[match(rownames(.), state.name)]`

## Excercise 9:

### Titanic Survival (from titanic package or datasets::Titanic)

**Dataset:** `as.data.frame(Titanic)`

**Task:**

- ▶ Use `dplyr` to calculate survival rate per class and sex.
- ▶ Use `ggplot2` to create a stacked bar plot showing survival by class and sex.